

CS158_HW4_____

November 30, 2020

```
[522]: import numpy as np
from numpy.linalg import norm
import math
from urllib.request import urlopen
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import linear_model
from urllib.request import urlopen
import random

import urllib
import scipy.optimize
import random
from collections import defaultdict # Dictionaries with default values
import nltk
from nltk.util import ngrams
import string
from nltk.stem.porter import PorterStemmer
import ast

import gzip
from collections import defaultdict
```

```
[ ]:
```

```
[406]: def parseDataFromURL(fname):
        for l in urlopen(fname):
            yield eval(l)

        def parseData(fname):
            for l in open(fname):
                yield eval(l)
```

```
[ ]: print("Reading data...")
data = list(parseData("./assignment1/train_Category.json"))
print("done")
```

```
df = pd.DataFrame(data)
df
```

```
[408]: training_data = data[0:10000]
training_data[0]
```

```
[408]: {'userID': 'u74382925',
        'genre': 'Adventure',
        'early_access': False,
        'reviewID': 'r75487422',
        'hours': 4.1,
        'text': 'Short Review:\nA good starting chapter for this series, despite the
main character being annoying (for now) and a short length. The story is good
and actually gets more interesting. Worth the try.\nLong Review:\nBlackwell
Legacy is the first on the series of (supposedly) 5 games that talks about the
main protagonist, Rosangela Blackwell, as being a so called Medium, and in this
first chapter we get to know how her story will start and how she will meet her
adventure companion Joey...and really, that\'s really all for for now and
that\'s not a bad thing, because in a way this game wants to show how hard her
new job is, and that she cannot escape her destiny as a medium.\nMy biggest
complain for this chapter, except the short length, it\'s the main protagonist
being a "bit" too annoying to be likeable, and most of her dialogues will always
be about complaining or just be annoyed. Understandable, sure, but lighten\' up
will ya!?\nHowever, considering that in the next installments she will be much
more likeable and kind of interesting, I\'d say give it a shot and see if you
like it: if you hate this first game, you might like the next, or can always
stop here.\nI recommend it.',
        'genreID': 3,
        'date': '2014-02-07'}
```

0.1 1.

```
[409]: BigramCount = defaultdict(int)
totalBigrams = 0

punct = string.punctuation

for d in training_data:
    t = d['text']
    t = t.lower() # lowercase string
    t = [c for c in t if not (c in punct)] # non-punct characters
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes
    for i in range(len(words)-1):
        totalBigrams += 1
```

```
BigramCount[(words[i], words[i+1])] += 1
```

```
[410]: counts = [(BigramCount[b], b) for b in BigramCount]
counts.sort()
counts.reverse()
print('Unique Bigrams:', len(counts))

print('\nTop 5 most frequent Bigrams')
counts[:5]
```

Unique Bigrams: 256618

Top 5 most frequent Bigrams

```
[410]: [(4441, ('this', 'game')),
(4249, ('the', 'game')),
(3359, ('of', 'the')),
(2020, ('if', 'you')),
(2017, ('in', 'the'))]
```

```
[ ]:
```

0.2 2.

```
[411]: bigrams = [b[1] for b in counts[:1000]]
bigramId = dict(zip(bigrams, range(len(bigrams))))
bigramSet = set(bigrams)
print(len(bigramSet))
```

1000

```
[412]: def feature(datum):
    feat = [0]*len(bigramSet)
    t = datum['text']
    t = t.lower() # lowercase string
    t = [c for c in t if not (c in punct)] # non-punct characters
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes
    for i in range(len(words)-1):
        b = (words[i], words[i+1])
        if not (b in bigramSet): continue
        feat[bigramId[b]] += 1
    feat.append(1)
    return feat
```

```
[413]: X = [feature(d) for d in training_data]
y = [np.log2(d['hours'] + 1) for d in training_data]
```

```
[414]: clf = linear_model.Ridge(1.0, fit_intercept=False) # MSE + 1.0 l2
        clf.fit(X, y)
        theta = clf.coef_
        predictions = clf.predict(X)
```

```
[415]: MSE = sum((predictions - y)**2) / len(y)
        print('MSE =', MSE)
```

MSE = 4.39378724720003

```
[ ]:
```

0.3 3.

```
[416]: # Get top 1000 unigrams
        unigramCount = defaultdict(int)
        totalUnigrams = 0

        for d in training_data:
            t = d['text']
            t = t.lower() # lowercase string
            t = [c for c in t if not (c in punct)] # non-punct characters
            t = ''.join(t) # convert back to string
            words = t.strip().split() # tokenizes
            for w in words:
                totalUnigrams += 1
                unigramCount[w] += 1
```

```
[417]: countsUni = [(unigramCount[w], w) for w in unigramCount]
        countsUni.sort()
        countsUni.reverse()
        print(len(countsUni))
```

29692

```
[418]: # Strategy: Combine top 500 unigrams and top 500 bigrams together
        unigrams = [w[1] for w in countsUni[:500]]
        unigramId = dict(zip(unigrams, range(len(unigrams))))
        unigramSet = set(unigrams)
        print(len(unigramSet))

        bigrams = [b[1] for b in counts[:500]]
        bigramId = dict(zip(bigrams, range(len(unigrams), len(unigrams) +
            → len(bigrams))))
        bigramSet = set(bigrams)
        print(len(bigramSet))
```

500
500

```
[419]: def feature2(datum):
    feat = [0]*(len(unigramSet) + len(bigramSet))
    t = datum['text']
    t = t.lower() # lowercase string
    t = [c for c in t if not (c in punct)] # non-punct characters
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes
    for w in words:
        if not (w in unigramSet): continue
        feat[unigramId[w]] += 1

    for i in range(len(words)-1):
        b = (words[i], words[i+1])
        if not (b in bigramSet): continue
        feat[bigramId[b]] += 1
    feat.append(1)
    return feat
```

```
[420]: X = [feature2(d) for d in training_data]
y = [np.log2(d['hours'] + 1) for d in training_data]
```

```
[421]: clf = linear_model.Ridge(1.0, fit_intercept=False) # MSE + 1.0 l2
clf.fit(X, y)
theta = clf.coef_
predictions = clf.predict(X)
```

```
[422]: MSE = sum((predictions - y)**2) / len(y)
print('MSE =', MSE)
```

MSE = 4.249894468206029

```
[ ]:
```

0.4 4.

```
[423]: def idf(word, data):
    counter = 0
    for datum in data:
        t = datum['text']
        t = t.lower() # lowercase string
        t = [c for c in t if not (c in punct)] # non-punct characters
        t = ''.join(t) # convert back to string
        words = t.strip().split() # tokenizes
        if word in words:
```

```

        counter += 1

    return np.log10(len(data) / counter)

```

```

[424]: idf_words = ['destiny', 'annoying', 'likeable', 'chapter', 'interesting']

for w in idf_words:
    idf_val = idf(w, training_data)
    print('IDF for ' + w + ' =', idf_val)

```

```

IDF for destiny = 3.3979400086720375
IDF for annoying = 1.8386319977650252
IDF for likeable = 3.0969100130080562
IDF for chapter = 2.221848749616356
IDF for interesting = 1.3585258894959005

```

```

[445]: # tf-idf of the five words in a document
def tf_idf(word, datum, data):
    counter = 0
    t = datum['text']
    t = t.lower() # lowercase string
    t = [c for c in t if not (c in punct)] # non-punct characters
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes

    for w in words:
        if w == word:
            counter += 1
    return counter * idf(word, data)

```

```

[426]: for d in training_data:
    if d['reviewID'] == 'r75487422':
        for w in idf_words:
            tf_idf_val = tf_idf(w, d, training_data)
            print('Tf-Idf for \'' + w + '\'' =', tf_idf_val)
        break

```

```

Tf-Idf for 'destiny' = 3.3979400086720375
Tf-Idf for 'annoying' = 3.6772639955300503
Tf-Idf for 'likeable' = 6.1938200260161125
Tf-Idf for 'chapter' = 6.665546248849068
Tf-Idf for 'interesting' = 2.717051778991801

```

```

[ ]:

```

0.5 5.

```
[511]: unigrams = [w[1] for w in countsUni[:1000]]
unigramId = dict(zip(unigrams, range(len(unigrams))))
unigramSet = set(unigrams)
print(len(unigramSet))

# Map to improve runtime of idf
def getIdfMap(unigramSet, data):
    idfCounter = defaultdict(float)
    for datum in data:
        t = datum['text']
        t = t.lower() # lowercase string
        t = [c for c in t if not (c in punct)] # non-punct characters
        t = ''.join(t) # convert back to string
        words = t.strip().split() # tokenizes
        for w in words:
            if w in unigramSet:
                idfCounter[w] += 1
    return idfCounter

idfCounter = getIdfMap(unigramSet, training_data)
```

1000

```
[512]: # tf-idf using faster implementation of idf
def tf_idf2(word, datum, data):
    counter = 0
    t = datum['text']
    t = t.lower() # lowercase string
    t = [c for c in t if not (c in punct)] # non-punct characters
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes

    for w in words:
        if w == word:
            counter += 1
    return counter * np.log10(len(data) / idfCounter[word])
```

```
[513]: # Unigrams only
def feature3(datum):
    feat = [0.0]*(len(unigramSet))
    t = datum['text']
    t = t.lower() # lowercase string
    t = [c for c in t if not (c in punct)] # non-punct characters
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes
    for w in words:
```

```

        if not (w in unigramSet): continue
        feat[unigramId[w]] = tf_idf2(w, datum, training_data)
    feat.append(1)
    return feat

```

```

[514]: X = [feature3(d) for d in training_data]
       y = [np.log2(d['hours'] + 1) for d in training_data]

```

```

[515]: clf = linear_model.Ridge(1.0, fit_intercept=False) # MSE + 1.0 l2
       clf.fit(X, y)
       theta = clf.coef_
       predictions = clf.predict(X)

```

```

[516]: MSE = sum((predictions - y)**2) / len(y)
       print('MSE =', MSE)

```

MSE = 4.199776205842035

```
[ ]:
```

0.6 6.

```

[538]: def CosineSimilarity(a, b):
       return np.dot(a,b) / (norm(a) * norm(b))

```

```

[539]: r75487422_data = []

       for datum in training_data:
           if datum['reviewID'] == 'r75487422':
               r75487422_data = datum
               break

       r75487422_tf_idf = feature3(r75487422_data)

       best_cos_sim = -99999
       best_datum = []

       for datum in training_data:
           if datum['reviewID'] == 'r75487422':
               continue

           tf_idf_rep = feature3(datum)
           cos_sim = CosineSimilarity(r75487422_tf_idf, tf_idf_rep)

           if best_cos_sim == -99999:
               best_cos_sim = cos_sim

```



```

        best_datum = datum
    elif cos_sim > best_cos_sim:
        best_cos_sim = cos_sim
        best_datum = datum

print('ReviewID = r75487422\n')
print('Review Text:\n', r75487422_data['text'])

print('\n\n')
print('Best Cosine Similarity compared to r75487422\n')
print('ReviewID =', best_datum['reviewID'])
print('Review Text:\n', best_datum['text'])
print('\n')
print('Cosine Similarity =', best_cos_sim)

```

ReviewID = r75487422

Review Text:

Short Review:

A good starting chapter for this series, despite the main character being annoying (for now) and a short length. The story is good and actually gets more interesting. Worth the try.

Long Review:

Blackwell Legacy is the first on the series of (supposedly) 5 games that talks about the main protagonist, Rosangela Blackwell, as being a so called Medium, and in this first chapter we get to know how her story will start and how she will meet her adventure companion Joey...and really, that's really all for for now and that's not a bad thing, because in a way this game wants to show how hard her new job is, and that she cannot escape her destiny as a medium.

My biggest complain for this chapter, except the short length, it's the main protagonist being a "bit" too annoying to be likeable, and most of her dialogues will always be about complaining or just be annoyed. Understandable, sure, but lighten' up will ya!?

However, considering that in the next installments she will be much more likeable and kind of interesting, I'd say give it a shot and see if you like it: if you hate this first game, you might like the next, or can always stop here. I recommend it.

Best Cosine Similarity compared to r75487422

ReviewID = r33215456

Review Text:

Before I tell you about RimWorld, let me tell you a story that happened in RimWorld. It's about that girl up there, drinking a beer. If you're not

convinced to jump in by the end of this tale, then we have nothing more to talk about. We can't be friends. Everyone else: we're still cool. So here it is, the story of Min, a pop star with a privileged upbringing, who is about to come crashing down to earth.

Min slept with all her animals around her. Since the crash, the 16-year-old pop idol was finding it hard to adjust. She would not help with heavy lifting, or farming, or pulling up weeds. The only thing she would do is take care of the colony's animals. Rabbits, dogs, tortoises. She loved them all. She even formed a close bond with a rat, whom she named 'Illness'.

But she wasn't alone. In this game, you can create your own scenario before crash-landing onto the planet's surface, adding or removing starter items (wood, steel, pelts, gold, widescreen TVs). You can also set the number of colonists and tweak dozens of other small rules, like compulsory character traits or rules about the planet's climate. The game let's you mould your scenario completely, right down to the flavour text.

As you can see, my five colonists had all escaped from a crashing party yacht. They were all rich, famous or both. A jeweler, a wealthy businessman, an ace fighter pilot, a star surgeon and (of course) Min the pop star. They all landed with severe hangovers. I am not kidding, this is a setting you can change.

Fast forward some weeks. The colonists had settled in. They had even survived their first few pirate raids. During one of these raids, they captured a large pirate with a low IQ called Ferdnand. This man was physically capable of nothing but sculpting statues and communing with animals. He was like Hodor with a paintbrush. FERDNAND! Eventually, we convinced him to join the colony.

I set him the task of taming a Megatherium - a huge and powerful ground sloth - so that the colony would have some extra muscle. But Min the pop idol was not having it. She had always been the colony's animal trainer, and she had little 'Illness' to prove it. Spurned, she decided to take on the job herself. She left the safety of the village walls and approached the huge beast with some corn in hand. It went mad.

Ferdnand was already on his way to train his new would-be pet when he saw Min being savaged. He took out his Uzi and started firing with surprising accuracy. FERDNAND! Every shot landed. But the Megatherium turned and charged. By the time Ferdnand started to run, it was too late. The animal ripped him to pieces and he died on the spot.

The other colonists formed up. They lured the beast into the walls, where they riddled it with bullets until dead. The colony's nudist surgeon (who weeks ago had experienced a mental breakdown and now did everything naked to keep himself happy) rushed out to rescue Min, who was lying in a bloody heap, writhing in unbearable pain. He took her to the clinic, wrapped her in bandages and made sure she got everything she needed - medicine, food and rest.

Well, maybe not rest.

You see, I had already scheduled the clinic's floor to be renovated during this time. And as Min slept (or tried to sleep) my other colonists were busy ripping up planks and replacing them with shiny white tiles, sending sparks and wood chips everywhere. The noise was unbearable. Min snapped.

Crazed with pain and lack of sleep, she went berserk and chased everyone out of the clinic. She terrorised the dinner guests in the common room, and pursued

them into the garden, where she chased them around. When she couldn't catch the colonists, she turned on the animals - hares, pigs and rats - punching the poor creatures to death in her rage, ignoring the scratches and bites she got in return. Ignoring her old friends.

Finally, a colonist stepped up, walking up to the rampaging diva and knocking her out cold. When Min came to, she was back in the clinic. It was completely refurbished, clean and presentable. But the same could not be said for Min. This once beautiful pop idol, who people would fall in love with before she even started to sing, was in tatters. She now had no nose, no left ear and no right arm.

Over the next two weeks she locked herself in her room and fell into terrible moods. She stopped handling the animals and was assigned to work only on sculptures. Her first work of art depicted a chicken vomiting on the shoes of a lawyer. I put it in the centre of her workshop.

One day, a ship appeared in orbit. They had items to trade. Among them was a prosthetic limb called a "power claw". After some thought, I decided it was better than nothing. And besides, these rich castaways could afford it. The nudist was scheduled to do the surgery. Min was getting her arm back, even if it was a claw.

When she woke this time, she went straight to her room and locked herself in. I wondered what was wrong. It wasn't until she was back at her post, carving new statues that I realised what had happened.

We had amputated the wrong arm.

Now, not only was Min a noseless, earless, armless monster, she also had a terrifying claw instead of her only remaining hand.

Welcome to RimWorld, where stories like this pointless tragedy are commonplace. It can probably best be summed up as: "the game for everyone who wanted to get into Dwarf Fortress but couldn't because ASCII." That's not say it doesn't have its own share of confusing menus. Supposedly there is a fluid teaching system but it seemed more to me like hints handed out at random. In the absence of a proper tutorial, it takes a lot of clicking and umming to understand why everything functions the way it does.

Let's look at the colonists, for instance. For the most part, they run around of their own volition but you can still take direct control in a number of ways. When I said that Ferdnand took out his gun, I meant that I clicked on Ferdnand and clicked the "draft" button, putting him into military mode - something you'll need to survive raids and attacks. But when I said he was already on his way to the gigantic animal, I meant he was heading to it autonomously, because it had been assigned to be "tamed" and he was a designated animal handler. There's a task grid that gives some order to it all. You can select which jobs you want each individual colonist to do and which jobs you want them to ignore. An advanced version lets you organise these even further by priority, fiddling with the numbers until you have each person specialised in a strict set of tasks. This is hidden away under "manual" control but it is something that soon becomes essential to learn.

It looks far more complicated than it is. And at any time you can just click on a colonist and right-click on something you want "prioritised" to have them commit to this instead of their job. It's a bit fiddly but comes in useful in a

pinch, for example, when someone insists on milking a cow in a burning barn instead of putting the fire out.

Building your settlement is a lot more straightforward. Prison Architect is the obvious inspiration for the way much of RimWorld's construction and furniture placement works. You drag lines of walls made of wood, steel or stone until you have what you want, filling it with nice floors, beds, billiards tables, lamps and so on. Power is supplied by solar panels, windmills, wood-burning generators. And electricity is stored in batteries and disseminated by carefully laid power conduits. Sometimes these power cables like to explode, sending your whole colony into darkness and spoiling all the food in your freezer. Sometimes the batteries get wet and go on fire. Sometimes that fire spreads to your sleeping quarters and burns everybody to death. How whimsical!

It's these accidents and mistakes that make it much more than just another management sim. And coupled with the finer details of the simulation, it makes for som

Cosine Similarity = 0.49042613059213225

[]:

0.7 7.

```
[706]: # Shuffle data
random.shuffle(data)

training_data = data[:10000]
validation_data = data[10000:20000]
test_data = data[20000:30000]

punct = string.punctuation

[650]: # Gets IDF Counts of Unigram (Remove or Preserve Punctuation)
def getIdfMapUnigram(unigramSet, data, rp):
    idfCounter = defaultdict(float)
    for datum in data:
        t = datum['text']
        t = t.lower() # lowercase string
        if rp == True:
            t = [c for c in t if not (c in punct)] # non-punct characters

        # Keep punctuation
        else:
            new_arr = []
            for c in t:
                if not (c in punct):
                    new_arr.append(c)
```

```

        else:
            new_arr.extend([' ', c, ' '])
        t = new_arr
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes
    for w in words:
        if w in unigramSet:
            idfCounter[w] += 1
    return idfCounter

# Gets IDF Counts of Bigram (Remove or Preserve Punctuation)
def getIdfMapBigram(bigramSet, data, rp):
    idfCounter = defaultdict(float)
    for datum in data:
        t = datum['text']
        t = t.lower() # lowercase string
        if rp == True:
            t = [c for c in t if not (c in punct)] # non-punct characters

        # Keep punctuation
        else:
            new_arr = []
            for c in t:
                if not (c in punct):
                    new_arr.append(c)
                else:
                    new_arr.extend([' ', c, ' '])
            t = new_arr
        t = ''.join(t) # convert back to string
        words = t.strip().split() # tokenizes
        for i in range(len(words)-1):
            b = (words[i], words[i+1])
            if b in bigramSet:
                idfCounter[b] += 1
    return idfCounter

# Calculate tf-idf scores for unigram or bigram along with punctuation remove
↳ or preserve
def tf_idf7(gram, datum, data, ub, rp, idfCounter):
    counter = 0
    t = datum['text']
    t = t.lower() # lowercase string
    if rp == True:
        t = [c for c in t if not (c in punct)] # non-punct characters

    # Keep punctuation
    else:

```

```

new_arr = []
for c in t:
    if not (c in punct):
        new_arr.append(c)
    else:
        new_arr.extend([' ', c, ' '])
t = new_arr
t = ''.join(t) # convert back to string
words = t.strip().split() # tokenizes

if ub == True:
    for w in words:
        if w == gram:
            counter += 1
else:
    for i in range(len(words)-1):
        b = (words[i], words[i+1])
        if b == gram:
            counter += 1
return counter * np.log10(len(data) / idfCounter[gram])

```

```

[651]: # Creates word feature for the feature matrix
def feature7(datum, Id, Set, ub, rp, tc, idfCounter):

    feat = [0]*(len(Set))

    t = datum['text']
    t = t.lower() # lowercase string
    # Remove punctuation
    if rp == True:
        t = [c for c in t if not (c in punct)] # non-punct characters

    # Keep punctuation
    else:
        new_arr = []
        for c in t:
            if not (c in punct):
                new_arr.append(c)
            else:
                new_arr.extend([' ', c, ' '])
        t = new_arr
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes

    # unigrams
    if ub == True:

```

```

    for w in words:
        if not (w in Set): continue

        # True = tfidf
        if tc == True:
            feat[Id[w]] = tf_idf7(w, datum, training_data, ub, rp,
↳idfCounter)
            # False = count
        else:
            feat[Id[w]] += 1

    # bigrams
    else:
        for i in range(len(words)-1):
            b = (words[i], words[i+1])
            if not (b in Set): continue
            # True = tfidf
            if tc == True:
                feat[Id[b]] = tf_idf7(b, datum, training_data, ub, rp,
↳idfCounter)
                # False = count
            else:
                feat[Id[b]] += 1

    feat.append(1)
    return feat

# Gets the dictionary set for top 1000 unigrams, bigrams with punctuation/no
↳punctuation
def q7Set(ub, rp, tc):

    # False = bigram
    if ub == False:
        BigramCount = defaultdict(int)
        totalBigrams = 0

    for d in training_data:
        t = d['text']
        t = t.lower() # lowercase string

        # Remove punctuation
        if rp == True:
            t = [c for c in t if not (c in punct)] # non-punct characters

        # Keep punctuation
        else:
            new_arr = []
            for c in t:

```

```

        if not (c in punct):
            new_arr.append(c)
        else:
            new_arr.extend([' ', c, ' '])
    t = new_arr
    t = ''.join(t) # convert back to string
    words = t.strip().split() # tokenizes
    for i in range(len(words)-1):
        totalBigrams += 1
        BigramCount[(words[i], words[i+1])] += 1

counts = [(BigramCount[b], b) for b in BigramCount]
counts.sort()
counts.reverse()

bigrams = [b[1] for b in counts[:1000]]
bigramId = dict(zip(bigrams, range(len(bigrams))))
bigramSet = set(bigrams)

return bigrams, bigramId, bigramSet

else:

    unigramCount = defaultdict(int)
    totalUnigrams = 0

    for d in training_data:
        t = d['text']
        t = t.lower() # lowercase string
        # Remove punctuation
        if rp == True:
            t = [c for c in t if not (c in punct)] # non-punct characters

        # Keep punctuation
        else:
            new_arr = []
            for c in t:
                if not (c in punct):
                    new_arr.append(c)
                else:
                    new_arr.extend([' ', c, ' '])
            t = new_arr
        t = ''.join(t) # convert back to string

```



```

        words = t.strip().split() # tokenizes
        for w in words:
            totalUnigrams += 1
            unigramCount[w] += 1

countsUni = [(unigramCount[w], w) for w in unigramCount]
countsUni.sort()
countsUni.reverse()

unigrams = [w[1] for w in countsUni[:1000]]
unigramId = dict(zip(unigrams, range(len(unigrams))))
unigramSet = set(unigrams)

return unigrams, unigramId, unigramSet

```

```

[707]: UB_Map = {True: 'Unigram', False: 'Bigram'}
        RP_Map = {True: 'Remove Punctuation', False: 'Preserve Punctuation'}
        TC_Map = {True: 'Tf-Idf Scores', False: 'Word Counts'}

Cs = [0.01, 0.1, 1, 10, 100]

# Lists to construct performance table
C_df = []

uni_df = []
bi_df = []

rp_df = []
pp_df = []

tfidf_df = []
wc_df = []

val_mse_df = []
test_mse_df = []

```

```

[708]: # Unigrams, Remove Punctuation, tfidf scores

grams, gramId, gramSet = q7Set(True, True, True)

IdfCounter = getIdfMapUnigram(gramSet, training_data, True)

X_train = [feature7(d, gramId, gramSet, True, True, True, IdfCounter) for d in
            ↪training_data]
y_train = [np.log2(d['hours'] + 1) for d in training_data]

```

```

X_validation = [feature7(d, gramId, gramSet, True, True, True, IdfCounter) for
    ↪ d in validation_data]
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]

X_test = [feature7(d, gramId, gramSet, True, True, True, IdfCounter) for d in
    ↪ test_data]
y_test = [np.log2(d['hours'] + 1) for d in test_data]

for C in Cs:
    clf = linear_model.Ridge(C, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_

    predictions = clf.predict(X_validation)
    MSE = sum((predictions - y_validation)**2) / len(y_validation)
    print('MSE Validation C =', C, '(' + UB_Map[True] + ', ' + RP_Map[True] + ', '
    ↪ ' + TC_Map[True] + ') =', MSE)

    val_mse_df.append(MSE)

    predictions = clf.predict(X_test)
    MSE = sum((predictions - y_test)**2) / len(y_test)
    print('MSE Test C =', C, '(' + UB_Map[True] + ', ' + RP_Map[True] + ', ' +
    ↪ TC_Map[True] + ') =', MSE)

    test_mse_df.append(MSE)

    print()

    C_df.append(C)
    uni_df.append(True)
    bi_df.append(False)

    rp_df.append(True)
    pp_df.append(False)

    tfidf_df.append(True)
    wc_df.append(False)

```

MSE Validation C = 0.01 (Unigram, Remove Punctuation, Tf-Idf Scores) =
5.673806372335936

MSE Test C = 0.01 (Unigram, Remove Punctuation, Tf-Idf Scores) =

5.44341643454938

MSE Validation C = 0.1 (Unigram, Remove Punctuation, Tf-Idf Scores) =
5.664963930672582

MSE Test C = 0.1 (Unigram, Remove Punctuation, Tf-Idf Scores) =
5.442522596772181

MSE Validation C = 1 (Unigram, Remove Punctuation, Tf-Idf Scores) =
5.644370081585309

MSE Test C = 1 (Unigram, Remove Punctuation, Tf-Idf Scores) = 5.442046784564783

MSE Validation C = 10 (Unigram, Remove Punctuation, Tf-Idf Scores) =
5.5789306327577215

MSE Test C = 10 (Unigram, Remove Punctuation, Tf-Idf Scores) = 5.397073383016831

MSE Validation C = 100 (Unigram, Remove Punctuation, Tf-Idf Scores) =
5.2129695143827535

MSE Test C = 100 (Unigram, Remove Punctuation, Tf-Idf Scores) =
5.122243582557432

```
[709]: # Bigrams, Remove punctation, tfidf scores
grams, gramId, gramSet = q7Set(False, True, True)

IdfCounter = getIdfMapBigram(gramSet, training_data, True)

X_train = [feature7(d, gramId, gramSet, False, True, True, IdfCounter) for d in
    ↪ training_data]
y_train = [np.log2(d['hours'] + 1) for d in training_data]

X_validation = [feature7(d, gramId, gramSet, False, True, True, IdfCounter) for
    ↪ d in validation_data]
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]

X_test = [feature7(d, gramId, gramSet, False, True, True, IdfCounter) for d in
    ↪ test_data]
y_test = [np.log2(d['hours'] + 1) for d in test_data]

for C in Cs:
    clf = linear_model.Ridge(C, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_

    predictions = clf.predict(X_validation)
    MSE = sum((predictions - y_validation)**2) / len(y_validation)
```

```

print('MSE Validation C =', C, '(' + UB_Map[False] + ', ' + RP_Map[True] +
↪', ' + TC_Map[True] + ') =', MSE)

val_mse_df.append(MSE)

predictions = clf.predict(X_test)
MSE = sum((predictions - y_test)**2) / len(y_test)
print('MSE Test C =', C, '(' + UB_Map[False] + ', ' + RP_Map[True] + ', ' +
↪TC_Map[True] + ') =', MSE)

test_mse_df.append(MSE)

print()

C_df.append(C)
uni_df.append(False)
bi_df.append(True)

rp_df.append(True)
pp_df.append(False)

tfidf_df.append(True)
wc_df.append(False)

```

MSE Validation C = 0.01 (Bigram, Remove Punctuation, Tf-Idf Scores) =
5.673287403837088

MSE Test C = 0.01 (Bigram, Remove Punctuation, Tf-Idf Scores) =
5.506209233571954

MSE Validation C = 0.1 (Bigram, Remove Punctuation, Tf-Idf Scores) =
5.672797087486668

MSE Test C = 0.1 (Bigram, Remove Punctuation, Tf-Idf Scores) = 5.505538566595185

MSE Validation C = 1 (Bigram, Remove Punctuation, Tf-Idf Scores) =
5.668004372267518

MSE Test C = 1 (Bigram, Remove Punctuation, Tf-Idf Scores) = 5.499240629174145

MSE Validation C = 10 (Bigram, Remove Punctuation, Tf-Idf Scores) =
5.621993571271137

MSE Test C = 10 (Bigram, Remove Punctuation, Tf-Idf Scores) = 5.448469226938605

MSE Validation C = 100 (Bigram, Remove Punctuation, Tf-Idf Scores) =
5.366705958094359

MSE Test C = 100 (Bigram, Remove Punctuation, Tf-Idf Scores) = 5.215462371009729

```

[710]: # Unigrams, Preserve Punctuation, tfidf scores
grams, gramId, gramSet = q7Set(True, False, True)

IdfCounter = getIdfMapUnigram(gramSet, training_data, False)

X_train = [feature7(d, gramId, gramSet, True, False, True, IdfCounter) for d in
    ↪training_data]
y_train = [np.log2(d['hours'] + 1) for d in training_data]

X_validation = [feature7(d, gramId, gramSet, True, False, True, IdfCounter) for
    ↪d in validation_data]
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]

X_test = [feature7(d, gramId, gramSet, True, False, True, IdfCounter) for d in
    ↪test_data]
y_test = [np.log2(d['hours'] + 1) for d in test_data]

for C in Cs:
    clf = linear_model.Ridge(C, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_

    predictions = clf.predict(X_validation)
    MSE = sum((predictions - y_validation)**2) / len(y_validation)
    print('MSE Validation C =', C, '(' + UB_Map[True] + ', ' + RP_Map[False] +
    ↪', ' + TC_Map[True] + ') =', MSE)

    val_mse_df.append(MSE)

    predictions = clf.predict(X_test)
    MSE = sum((predictions - y_test)**2) / len(y_test)
    print('MSE Test C =', C, '(' + UB_Map[True] + ', ' + RP_Map[False] + ', ' +
    ↪TC_Map[True] + ') =', MSE)

    test_mse_df.append(MSE)

    print()

    C_df.append(C)
    uni_df.append(True)
    bi_df.append(False)

    rp_df.append(False)
    pp_df.append(True)

    tfidf_df.append(True)

```

```
wc_df.append(False)
```

```
MSE Validation C = 0.01 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.570218397178122
```

```
MSE Test C = 0.01 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.422839226112137
```

```
MSE Validation C = 0.1 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.568655202923451
```

```
MSE Test C = 0.1 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.4211872568211925
```

```
MSE Validation C = 1 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.567519236384305
```

```
MSE Test C = 1 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.419563281432396
```

```
MSE Validation C = 10 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.517231538305895
```

```
MSE Test C = 10 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.376351913572079
```

```
MSE Validation C = 100 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.202233460893193
```

```
MSE Test C = 100 (Unigram, Preserve Punctuation, Tf-Idf Scores) =  
5.11910178254146
```

```
[711]: # Bigrams, Preserve Punctuation, tfidf scores  
grams, gramId, gramSet = q7Set(False, False, True)  
  
IdfCounter = getIdfMapBigram(gramSet, training_data, False)  
  
X_train = [feature7(d, gramId, gramSet, False, False, True, IdfCounter) for d  
            ↪ in training_data]  
y_train = [np.log2(d['hours'] + 1) for d in training_data]  
  
X_validation = [feature7(d, gramId, gramSet, False, False, True, IdfCounter)  
                ↪ for d in validation_data]  
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]  
  
X_test = [feature7(d, gramId, gramSet, False, False, True, IdfCounter) for d in  
           ↪ test_data]  
y_test = [np.log2(d['hours'] + 1) for d in test_data]  
  
for C in Cs:
```

```

clf = linear_model.Ridge(C, fit_intercept=False)
clf.fit(X_train, y_train)
theta = clf.coef_

predictions = clf.predict(X_validation)
MSE = sum((predictions - y_validation)**2) / len(y_validation)
print('MSE Validation C =', C, '(' + UB_Map[False] + ', ' + RP_Map[False] + '
→', ' + TC_Map[True] + ') =', MSE)

val_mse_df.append(MSE)

predictions = clf.predict(X_test)
MSE = sum((predictions - y_test)**2) / len(y_test)
print('MSE Test C =', C, '(' + UB_Map[False] + ', ' + RP_Map[False] + ', ' +
→TC_Map[True] + ') =', MSE)

test_mse_df.append(MSE)

print()

C_df.append(C)
uni_df.append(False)
bi_df.append(True)

rp_df.append(False)
pp_df.append(True)

tfidf_df.append(True)
wc_df.append(False)

```

MSE Validation C = 0.01 (Bigram, Preserve Punctuation, Tf-Idf Scores) =
5.582117883160818

MSE Test C = 0.01 (Bigram, Preserve Punctuation, Tf-Idf Scores) =
5.525768841253606

MSE Validation C = 0.1 (Bigram, Preserve Punctuation, Tf-Idf Scores) =
5.5811412448533835

MSE Test C = 0.1 (Bigram, Preserve Punctuation, Tf-Idf Scores) =
5.524634175402464

MSE Validation C = 1 (Bigram, Preserve Punctuation, Tf-Idf Scores) =
5.574056215986605

MSE Test C = 1 (Bigram, Preserve Punctuation, Tf-Idf Scores) = 5.515303428263916

MSE Validation C = 10 (Bigram, Preserve Punctuation, Tf-Idf Scores) =
5.53675784403914

MSE Test C = 10 (Bigram, Preserve Punctuation, Tf-Idf Scores) =

5.461732251629773

MSE Validation C = 100 (Bigram, Preserve Punctuation, Tf-Idf Scores) =

5.272388170686014

MSE Test C = 100 (Bigram, Preserve Punctuation, Tf-Idf Scores) =

5.232844454976385

```
[712]: # Unigrams, Remove Punctuation, word counts
grams, gramId, gramSet = q7Set(True, True, False)

IdfCounter = getIdfMapUnigram(gramSet, training_data, True)

X_train = [feature7(d, gramId, gramSet, True, True, False, IdfCounter) for d in
    ↪ training_data]
y_train = [np.log2(d['hours'] + 1) for d in training_data]

X_validation = [feature7(d, gramId, gramSet, True, True, False, IdfCounter) for
    ↪ d in validation_data]
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]

X_test = [feature7(d, gramId, gramSet, True, True, False, IdfCounter) for d in
    ↪ test_data]
y_test = [np.log2(d['hours'] + 1) for d in test_data]

for C in Cs:
    clf = linear_model.Ridge(C, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_

    predictions = clf.predict(X_validation)
    MSE = sum((predictions - y_validation)**2) / len(y_validation)
    print('MSE Validation C =', C, '(' + UB_Map[True] + ', ' + RP_Map[True] + ', '
    ↪ ' + TC_Map[False] + ') =', MSE)

    val_mse_df.append(MSE)

    predictions = clf.predict(X_test)
    MSE = sum((predictions - y_test)**2) / len(y_test)
    print('MSE Test C =', C, '(' + UB_Map[True] + ', ' + RP_Map[True] + ', ' +
    ↪ TC_Map[False] + ') =', MSE)

    test_mse_df.append(MSE)

    print()
```



```

C_df.append(C)
uni_df.append(True)
bi_df.append(False)
rp_df.append(True)
pp_df.append(False)
tfidf_df.append(False)
wc_df.append(True)

```

```

MSE Validation C = 0.01 (Unigram, Remove Punctuation, Word Counts) =
5.674954707791254
MSE Test C = 0.01 (Unigram, Remove Punctuation, Word Counts) = 5.443446979850613

MSE Validation C = 0.1 (Unigram, Remove Punctuation, Word Counts) =
5.672776933245138
MSE Test C = 0.1 (Unigram, Remove Punctuation, Word Counts) = 5.44173383018089

MSE Validation C = 1 (Unigram, Remove Punctuation, Word Counts) =
5.65154026758534
MSE Test C = 1 (Unigram, Remove Punctuation, Word Counts) = 5.425086121336876

MSE Validation C = 10 (Unigram, Remove Punctuation, Word Counts) =
5.481431537033336
MSE Test C = 10 (Unigram, Remove Punctuation, Word Counts) = 5.295566510408608

MSE Validation C = 100 (Unigram, Remove Punctuation, Word Counts) =
4.968279141466649
MSE Test C = 100 (Unigram, Remove Punctuation, Word Counts) = 4.949465612094928

```

```

[713]: # Bigrams, Remove Punctuation, word counts
grams, gramId, gramSet = q7Set(False, True, False)

IdfCounter = getIdfMapBigram(gramSet, training_data, True)

X_train = [feature7(d, gramId, gramSet, False, True, False, IdfCounter) for d
↳ in training_data]
y_train = [np.log2(d['hours'] + 1) for d in training_data]

X_validation = [feature7(d, gramId, gramSet, False, True, False, IdfCounter)
↳ for d in validation_data]
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]

X_test = [feature7(d, gramId, gramSet, False, True, False, IdfCounter) for d in
↳ test_data]
y_test = [np.log2(d['hours'] + 1) for d in test_data]

```

```

for C in Cs:
    clf = linear_model.Ridge(C, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_

    predictions = clf.predict(X_validation)
    MSE = sum((predictions - y_validation)**2) / len(y_validation)
    print('MSE Validation C =', C, '(' + UB_Map[False] + ', ' + RP_Map[True] + '
→', ' + TC_Map[False] + ') =', MSE)

    val_mse_df.append(MSE)

    predictions = clf.predict(X_test)
    MSE = sum((predictions - y_test)**2) / len(y_test)
    print('MSE Test C =', C, '(' + UB_Map[False] + ', ' + RP_Map[True] + ', ' +
→TC_Map[False] + ') =', MSE)

    test_mse_df.append(MSE)

    print()

    C_df.append(C)
    uni_df.append(False)
    bi_df.append(True)

    rp_df.append(True)
    pp_df.append(False)

    tfidf_df.append(False)
    wc_df.append(True)

```

MSE Validation C = 0.01 (Bigram, Remove Punctuation, Word Counts) =
5.67317280728575

MSE Test C = 0.01 (Bigram, Remove Punctuation, Word Counts) = 5.505991995034113

MSE Validation C = 0.1 (Bigram, Remove Punctuation, Word Counts) =
5.671681325319052

MSE Test C = 0.1 (Bigram, Remove Punctuation, Word Counts) = 5.503457637327854

MSE Validation C = 1 (Bigram, Remove Punctuation, Word Counts) =
5.657019481153744

MSE Test C = 1 (Bigram, Remove Punctuation, Word Counts) = 5.481632097339202

MSE Validation C = 10 (Bigram, Remove Punctuation, Word Counts) =
5.541543954145307

MSE Test C = 10 (Bigram, Remove Punctuation, Word Counts) = 5.344781053758389

MSE Validation C = 100 (Bigram, Remove Punctuation, Word Counts) =
5.287558047852009
MSE Test C = 100 (Bigram, Remove Punctuation, Word Counts) = 5.08111846367742

```
[714]: # Unigrams, Preserve Punctuation, word counts
grams, gramId, gramSet = q7Set(True, False, False)

IdfCounter = getIdfMapUnigram(gramSet, training_data, False)

X_train = [feature7(d, gramId, gramSet, True, False, False, IdfCounter) for d
    ↪in training_data]
y_train = [np.log2(d['hours'] + 1) for d in training_data]

X_validation = [feature7(d, gramId, gramSet, True, False, False, IdfCounter)
    ↪for d in validation_data]
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]

X_test = [feature7(d, gramId, gramSet, True, False, False, IdfCounter) for d in
    ↪test_data]
y_test = [np.log2(d['hours'] + 1) for d in test_data]

for C in Cs:
    clf = linear_model.Ridge(C, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_

    predictions = clf.predict(X_validation)
    MSE = sum((predictions - y_validation)**2) / len(y_validation)
    print('MSE Validation C =', C, '(' + UB_Map[True] + ', ' + RP_Map[False] +
    ↪', ' + TC_Map[False] + ') =', MSE)

    val_mse_df.append(MSE)

    predictions = clf.predict(X_test)
    MSE = sum((predictions - y_test)**2) / len(y_test)
    print('MSE Test C =', C, '(' + UB_Map[True] + ', ' + RP_Map[False] + ', ' +
    ↪TC_Map[False] + ') =', MSE)

    test_mse_df.append(MSE)

    print()

    C_df.append(C)
    uni_df.append(True)
    bi_df.append(False)
```

```
rp_df.append(False)
pp_df.append(True)

tfidf_df.append(False)
wc_df.append(True)
```

MSE Validation C = 0.01 (Unigram, Preserve Punctuation, Word Counts) =
5.570524252790048

MSE Test C = 0.01 (Unigram, Preserve Punctuation, Word Counts) =
5.423149116056916

MSE Validation C = 0.1 (Unigram, Preserve Punctuation, Word Counts) =
5.568800507111982

MSE Test C = 0.1 (Unigram, Preserve Punctuation, Word Counts) =
5.421671333791256

MSE Validation C = 1 (Unigram, Preserve Punctuation, Word Counts) =
5.551952768534205

MSE Test C = 1 (Unigram, Preserve Punctuation, Word Counts) = 5.407268381152867

MSE Validation C = 10 (Unigram, Preserve Punctuation, Word Counts) =
5.41471873855819

MSE Test C = 10 (Unigram, Preserve Punctuation, Word Counts) = 5.292599618963542

MSE Validation C = 100 (Unigram, Preserve Punctuation, Word Counts) =
4.978360758164593

MSE Test C = 100 (Unigram, Preserve Punctuation, Word Counts) = 4.96190152372311

```
[715]: # Bigrams, Preserve Punctuation, word counts
grams, gramId, gramSet = q7Set(False, False, False)

IdfCounter = getIdfMapBigram(gramSet, training_data, False)

X_train = [feature7(d, gramId, gramSet, False, False, False, IdfCounter) for d
    ↪ in training_data]
y_train = [np.log2(d['hours'] + 1) for d in training_data]

X_validation = [feature7(d, gramId, gramSet, False, False, False, IdfCounter)
    ↪ for d in validation_data]
y_validation = [np.log2(d['hours'] + 1) for d in validation_data]

X_test = [feature7(d, gramId, gramSet, False, False, False, IdfCounter) for d
    ↪ in test_data]
y_test = [np.log2(d['hours'] + 1) for d in test_data]
```

```

for C in Cs:
    clf = linear_model.Ridge(C, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_

    predictions = clf.predict(X_validation)
    MSE = sum((predictions - y_validation)**2) / len(y_validation)
    print('MSE Validation C =', C, '(' + UB_Map[False] + ', ' + RP_Map[False] + '
    ↪', ' + TC_Map[False] + ') =', MSE)

    val_mse_df.append(MSE)

    predictions = clf.predict(X_test)
    MSE = sum((predictions - y_test)**2) / len(y_test)
    print('MSE Test C =', C, '(' + UB_Map[False] + ', ' + RP_Map[False] + ', ' +
    ↪TC_Map[False] + ') =', MSE)

    test_mse_df.append(MSE)

    print()

    C_df.append(C)
    uni_df.append(False)
    bi_df.append(True)

    rp_df.append(False)
    pp_df.append(True)

    tfidf_df.append(False)
    wc_df.append(True)

```

MSE Validation C = 0.01 (Bigram, Preserve Punctuation, Word Counts) =
5.581761417866405

MSE Test C = 0.01 (Bigram, Preserve Punctuation, Word Counts) =
5.525502739262317

MSE Validation C = 0.1 (Bigram, Preserve Punctuation, Word Counts) =
5.5779367383972795

MSE Test C = 0.1 (Bigram, Preserve Punctuation, Word Counts) = 5.522169218695633

MSE Validation C = 1 (Bigram, Preserve Punctuation, Word Counts) =
5.553447642596715

MSE Test C = 1 (Bigram, Preserve Punctuation, Word Counts) = 5.498788372251529

MSE Validation C = 10 (Bigram, Preserve Punctuation, Word Counts) =
5.40154892954481
MSE Test C = 10 (Bigram, Preserve Punctuation, Word Counts) = 5.375797766715348

MSE Validation C = 100 (Bigram, Preserve Punctuation, Word Counts) =
5.064762706947522
MSE Test C = 100 (Bigram, Preserve Punctuation, Word Counts) = 5.103291923113969

```
[718]: des_df = []
for i in range(40):
    s = ""
    if uni_df[i]:
        s += "Unigram, "
    else:
        s += "Bigram, "

    if rp_df[i]:
        s += "Remove Punctuation, "
    else:
        s += "Preserve Punctuation, "

    if tfidf_df[i]:
        s += "Tf-Idf Scores"
    else:
        s += "Word Counts"

    des_df.append(s)

d = {
    'Model Description': des_df,
    'Regularization Parameter': C_df,
    'Validation MSE': val_mse_df,
    'Test MSE': test_mse_df
}
model_num = ['Model ' + str(n) for n in range(1,41)]
performance_table = pd.DataFrame(data=d, index=model_num, columns=['Model_
→Description', 'Regularization Parameter', 'Validation MSE', 'Test MSE'])
print('Performance Table')
performance_table
```

Performance Table

```
[718]:
```

	Model Description \
Model 1	Unigram, Remove Punctuation, Tf-Idf Scores
Model 2	Unigram, Remove Punctuation, Tf-Idf Scores

Model 3	Unigram, Remove Punctuation, Tf-Idf Scores
Model 4	Unigram, Remove Punctuation, Tf-Idf Scores
Model 5	Unigram, Remove Punctuation, Tf-Idf Scores
Model 6	Bigram, Remove Punctuation, Tf-Idf Scores
Model 7	Bigram, Remove Punctuation, Tf-Idf Scores
Model 8	Bigram, Remove Punctuation, Tf-Idf Scores
Model 9	Bigram, Remove Punctuation, Tf-Idf Scores
Model 10	Bigram, Remove Punctuation, Tf-Idf Scores
Model 11	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 12	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 13	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 14	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 15	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 16	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 17	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 18	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 19	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 20	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 21	Unigram, Remove Punctuation, Word Counts
Model 22	Unigram, Remove Punctuation, Word Counts
Model 23	Unigram, Remove Punctuation, Word Counts
Model 24	Unigram, Remove Punctuation, Word Counts
Model 25	Unigram, Remove Punctuation, Word Counts
Model 26	Bigram, Remove Punctuation, Word Counts
Model 27	Bigram, Remove Punctuation, Word Counts
Model 28	Bigram, Remove Punctuation, Word Counts
Model 29	Bigram, Remove Punctuation, Word Counts
Model 30	Bigram, Remove Punctuation, Word Counts
Model 31	Unigram, Preserve Punctuation, Word Counts
Model 32	Unigram, Preserve Punctuation, Word Counts
Model 33	Unigram, Preserve Punctuation, Word Counts
Model 34	Unigram, Preserve Punctuation, Word Counts
Model 35	Unigram, Preserve Punctuation, Word Counts
Model 36	Bigram, Preserve Punctuation, Word Counts
Model 37	Bigram, Preserve Punctuation, Word Counts
Model 38	Bigram, Preserve Punctuation, Word Counts
Model 39	Bigram, Preserve Punctuation, Word Counts
Model 40	Bigram, Preserve Punctuation, Word Counts

	Regularization Parameter	Validation MSE	Test MSE
Model 1	0.01	5.673806	5.443416
Model 2	0.10	5.664964	5.442523
Model 3	1.00	5.644370	5.442047
Model 4	10.00	5.578931	5.397073
Model 5	100.00	5.212970	5.122244
Model 6	0.01	5.673287	5.506209
Model 7	0.10	5.672797	5.505539

Model 8	1.00	5.668004	5.499241
Model 9	10.00	5.621994	5.448469
Model 10	100.00	5.366706	5.215462
Model 11	0.01	5.570218	5.422839
Model 12	0.10	5.568655	5.421187
Model 13	1.00	5.567519	5.419563
Model 14	10.00	5.517232	5.376352
Model 15	100.00	5.202233	5.119102
Model 16	0.01	5.582118	5.525769
Model 17	0.10	5.581141	5.524634
Model 18	1.00	5.574056	5.515303
Model 19	10.00	5.536758	5.461732
Model 20	100.00	5.272388	5.232844
Model 21	0.01	5.674955	5.443447
Model 22	0.10	5.672777	5.441734
Model 23	1.00	5.651540	5.425086
Model 24	10.00	5.481432	5.295567
Model 25	100.00	4.968279	4.949466
Model 26	0.01	5.673173	5.505992
Model 27	0.10	5.671681	5.503458
Model 28	1.00	5.657019	5.481632
Model 29	10.00	5.541544	5.344781
Model 30	100.00	5.287558	5.081118
Model 31	0.01	5.570524	5.423149
Model 32	0.10	5.568801	5.421671
Model 33	1.00	5.551953	5.407268
Model 34	10.00	5.414719	5.292600
Model 35	100.00	4.978361	4.961902
Model 36	0.01	5.581761	5.525503
Model 37	0.10	5.577937	5.522169
Model 38	1.00	5.553448	5.498788
Model 39	10.00	5.401549	5.375798
Model 40	100.00	5.064763	5.103292

[]:

```
[719]: print('Performance Table sorted by Validation MSE')
performance_table.sort_values('Validation MSE')
```

Performance Table sorted by Validation MSE

```
[719]:
```

	Model Description \
Model 25	Unigram, Remove Punctuation, Word Counts
Model 35	Unigram, Preserve Punctuation, Word Counts
Model 40	Bigram, Preserve Punctuation, Word Counts
Model 15	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 5	Unigram, Remove Punctuation, Tf-Idf Scores

Model 20	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 30	Bigram, Remove Punctuation, Word Counts
Model 10	Bigram, Remove Punctuation, Tf-Idf Scores
Model 39	Bigram, Preserve Punctuation, Word Counts
Model 34	Unigram, Preserve Punctuation, Word Counts
Model 24	Unigram, Remove Punctuation, Word Counts
Model 14	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 19	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 29	Bigram, Remove Punctuation, Word Counts
Model 33	Unigram, Preserve Punctuation, Word Counts
Model 38	Bigram, Preserve Punctuation, Word Counts
Model 13	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 12	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 32	Unigram, Preserve Punctuation, Word Counts
Model 11	Unigram, Preserve Punctuation, Tf-Idf Scores
Model 31	Unigram, Preserve Punctuation, Word Counts
Model 18	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 37	Bigram, Preserve Punctuation, Word Counts
Model 4	Unigram, Remove Punctuation, Tf-Idf Scores
Model 17	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 36	Bigram, Preserve Punctuation, Word Counts
Model 16	Bigram, Preserve Punctuation, Tf-Idf Scores
Model 9	Bigram, Remove Punctuation, Tf-Idf Scores
Model 3	Unigram, Remove Punctuation, Tf-Idf Scores
Model 23	Unigram, Remove Punctuation, Word Counts
Model 28	Bigram, Remove Punctuation, Word Counts
Model 2	Unigram, Remove Punctuation, Tf-Idf Scores
Model 8	Bigram, Remove Punctuation, Tf-Idf Scores
Model 27	Bigram, Remove Punctuation, Word Counts
Model 22	Unigram, Remove Punctuation, Word Counts
Model 7	Bigram, Remove Punctuation, Tf-Idf Scores
Model 26	Bigram, Remove Punctuation, Word Counts
Model 6	Bigram, Remove Punctuation, Tf-Idf Scores
Model 1	Unigram, Remove Punctuation, Tf-Idf Scores
Model 21	Unigram, Remove Punctuation, Word Counts

	Regularization Parameter	Validation MSE	Test MSE
Model 25	100.00	4.968279	4.949466
Model 35	100.00	4.978361	4.961902
Model 40	100.00	5.064763	5.103292
Model 15	100.00	5.202233	5.119102
Model 5	100.00	5.212970	5.122244
Model 20	100.00	5.272388	5.232844
Model 30	100.00	5.287558	5.081118
Model 10	100.00	5.366706	5.215462
Model 39	10.00	5.401549	5.375798
Model 34	10.00	5.414719	5.292600

Model 24	10.00	5.481432	5.295567
Model 14	10.00	5.517232	5.376352
Model 19	10.00	5.536758	5.461732
Model 29	10.00	5.541544	5.344781
Model 33	1.00	5.551953	5.407268
Model 38	1.00	5.553448	5.498788
Model 13	1.00	5.567519	5.419563
Model 12	0.10	5.568655	5.421187
Model 32	0.10	5.568801	5.421671
Model 11	0.01	5.570218	5.422839
Model 31	0.01	5.570524	5.423149
Model 18	1.00	5.574056	5.515303
Model 37	0.10	5.577937	5.522169
Model 4	10.00	5.578931	5.397073
Model 17	0.10	5.581141	5.524634
Model 36	0.01	5.581761	5.525503
Model 16	0.01	5.582118	5.525769
Model 9	10.00	5.621994	5.448469
Model 3	1.00	5.644370	5.442047
Model 23	1.00	5.651540	5.425086
Model 28	1.00	5.657019	5.481632
Model 2	0.10	5.664964	5.442523
Model 8	1.00	5.668004	5.499241
Model 27	0.10	5.671681	5.503458
Model 22	0.10	5.672777	5.441734
Model 7	0.10	5.672797	5.505539
Model 26	0.01	5.673173	5.505992
Model 6	0.01	5.673287	5.506209
Model 1	0.01	5.673806	5.443416
Model 21	0.01	5.674955	5.443447

After sorting the performance table by Validation MSE, the best performing model for this randomized train, validation, test split was Model 25 (Unigram, Remove Punctuation, Word Counts) because it had the smallest validation MSE compared to the other models. From the table we can also see that the top 8 models with the lowest validation MSEs are all of the different variation types (8 variations) where the regularization parameter was 100.

From our best performing model, we can see that the corresponding Test MSE for the best model (Model 25) is 4.949466

[]:

[]:

[]: